

PROJECT REPORT

NLP-BASED RESEARCH PAPER

RECOMMENDATION SYSTEM

1. Introduction

The vast and ever-expanding corpus of research literature poses a significant challenge for researchers who seek relevant academic papers quickly. With thousands of papers being published daily on repositories like arXiv, IEEE Xplore, and Springer, manual search becomes time-consuming and inefficient. To address this, the project focuses on building a hybrid research paper recommendation system that combines Natural Language Processing (NLP) and Information Retrieval (IR) techniques.

The system intelligently analyzes the content of a research paper (PDF input) and recommends similar or related papers from a large dataset using semantic similarity measures. This approach leverages sentence embeddings, latent semantic analysis (LSA), and FAISS-based vector search for high-speed and context-aware retrieval.

Objectives

- Automate the discovery of relevant academic literature.
- Enhance research efficiency using hybrid recommendation models.
- Combine both statistical and deep learning techniques for accurate semantic matching.
- Demonstrate real-world deployment using Google Colab and open-access datasets.

2. Literature Review

Traditional recommendation systems relied heavily on keyword matching and metadata such as titles, authors, and abstracts. While useful, such systems fail to capture deep semantic relationships between research topics.

Recent developments in Natural Language Processing and Representation Learning have improved semantic search significantly. Methods such as:

- TF-IDF (Term Frequency–Inverse Document Frequency) provide weighted keyword matching.
- Latent Semantic Analysis (LSA) using Singular Value Decomposition (SVD) uncovers hidden relationships in textual data.
- Sentence Transformers (e.g., all-MiniLM-L6-v2) encode entire sentences into dense numerical vectors that capture contextual meaning.
- FAISS (Facebook AI Similarity Search) enables fast similarity computation over millions of embeddings.

Hybrid approaches that integrate statistical models with deep embeddings have shown the best balance between accuracy, scalability, and speed, making them ideal for this system.

3. Methodology

3.1 Dataset Collection

The project uses the arXiv dataset from Cornell University, which contains millions of metadata records of scientific publications.

Using the **kagglehub** API, a subset of 50,000 papers was downloaded for experimental use.

Each record includes:

- title
- abstract
- Combined text: **title + abstract**

The data is stored in a pandas DataFrame for further preprocessing and modeling.

3.2 Data Preprocessing

Textual data is cleaned by:

- Removing stop words.
- Lowercasing.
- Tokenization for TF-IDF vectorization.

The combined text fields form the input corpus for both SVD and embedding models.

3.3 Model Architecture

The system is designed as a two-phase hybrid recommendation pipeline:

Phase 1: SVD/LSA-based Shortlisting

- A TF-IDF Vectorizer transforms the corpus into a high-dimensional sparse matrix (10,000 features).
- Truncated SVD ($n_components = 100$) reduces dimensionality while retaining semantic structure.
- Cosine similarity is computed between the input paper and the dataset to shortlist the top 100 most similar papers.

Mathematical representation:

$$TFIDF(t, d) = TF(t, d) \times \log \frac{N}{DF(t)}$$

$$\text{Similarity}(A, B) = \frac{A \cdot B}{||A|| \times ||B||}$$

Phase 2: Embedding-based Reranking

- The Sentence-BERT model (all-MiniLM-L6-v2) encodes text into 384-dimensional embeddings.
- Embeddings are L2 normalized and stored in a FAISS index for efficient inner product search.
- The top shortlisted papers are reranked based on embedding similarity scores.

$$\text{cosine_sim}(x, y) = \frac{x \cdot y}{||x|| \times ||y||}$$

3.4 PDF Text Extraction

The system supports user-uploaded research papers in PDF format. Using PyMuPDF (fitz):

- Each page is read sequentially.
- Extracted text is concatenated for downstream processing.

```
def extract_text_from_pdf(pdf_path):  
    doc = fitz.open(pdf_path)  
    text = ""  
    for page in doc:  
        text += page.get_text()  
    return text
```

3.5 Recommendation Pipeline

When a user uploads a PDF:


1. The text is extracted.
2. The TF-IDF + SVD model shortlists top 100 papers.
3. The Sentence-BERT model embeds both input and shortlisted papers.
4. FAISS computes the top 10 semantically similar papers.
5. The system displays results with title, similarity score, and abstract snippet.

4. Experimental Results

4.1 Model Performance

- SVD dimensionality: 100 components
- Embedding dimension: 384
- Dataset size: 50,001 papers
- R-squared (linear regression test): 0.9796 (used to validate model pipeline functionality)


4.2 Sample Output


 Please upload a research paper PDF...

 Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.


Saving Brain MRI Superresolution.pdf to Brain MRI Superresolution (1).pdf

 Top 10 Recommended Papers (Hybrid: SVD shortlist + Embedding rerank):

1.  Title: Automated identification of neurons and their locations


Score: 0.3386

Abstract: Individual locations of many neuronal cell bodies ($>10^4$) are needed to enable statistically significant measurements of spatial organization within the brain such as nearest-neighbor and microcolumnarity measurements. In this paper, we introduce an Automated Neuron Recognition Algorithm (ANRA) wh...

2.  Title: Automated detection of lung nodules in low-dose computed tomography


Score: 0.3368

Abstract: A computer-aided detection (CAD) system for the identification of pulmonary nodules in low-dose multi-detector computed-tomography (CT) images has been developed in the framework of the MAGIC-5 Italian project. One of the main goals of this project is to build a distributed database of lung CT sca...

3.  Title: An automated system for lung nodule detection in low-dose computed tomography


Score: 0.3332

Abstract: A computer-aided detection (CAD) system for the identification of pulmonary nodules in low-dose multi-detector helical Computed Tomography (CT) images was developed in the framework of the MAGIC-5 Italian project. One of the main goals of this project is to build a distributed database of lung CT ...

4.  Title: A Semi-parametric Technique for the Quantitative Analysis of Dynamic Contrast-enhanced MR Images Based on Bayesian P-splines

Score: 0.2965

Abstract: Dynamic Contrast-enhanced Magnetic Resonance Imaging (DCE-MRI) is an important tool for detecting subtle kinetic changes in cancerous tissue. Quantitative analysis of DCE-MRI typically involves the convolution of an arterial input function (AIF) with a nonlinear pharmacokinetic model of the contra...

5.  Title: Fuzzy Modeling of Electrical Impedance Tomography Image of the Lungs

Score: 0.2687

Abstract: Electrical Impedance Tomography (EIT) is a functional imaging method that is being developed for bedside use in critical care medicine. Aiming at improving the chest anatomical resolution of EIT images we developed a fuzzy model based on EIT high temporal resolution and the functional information ...

The results demonstrate that the system successfully retrieves contextually related research papers with high semantic similarity.

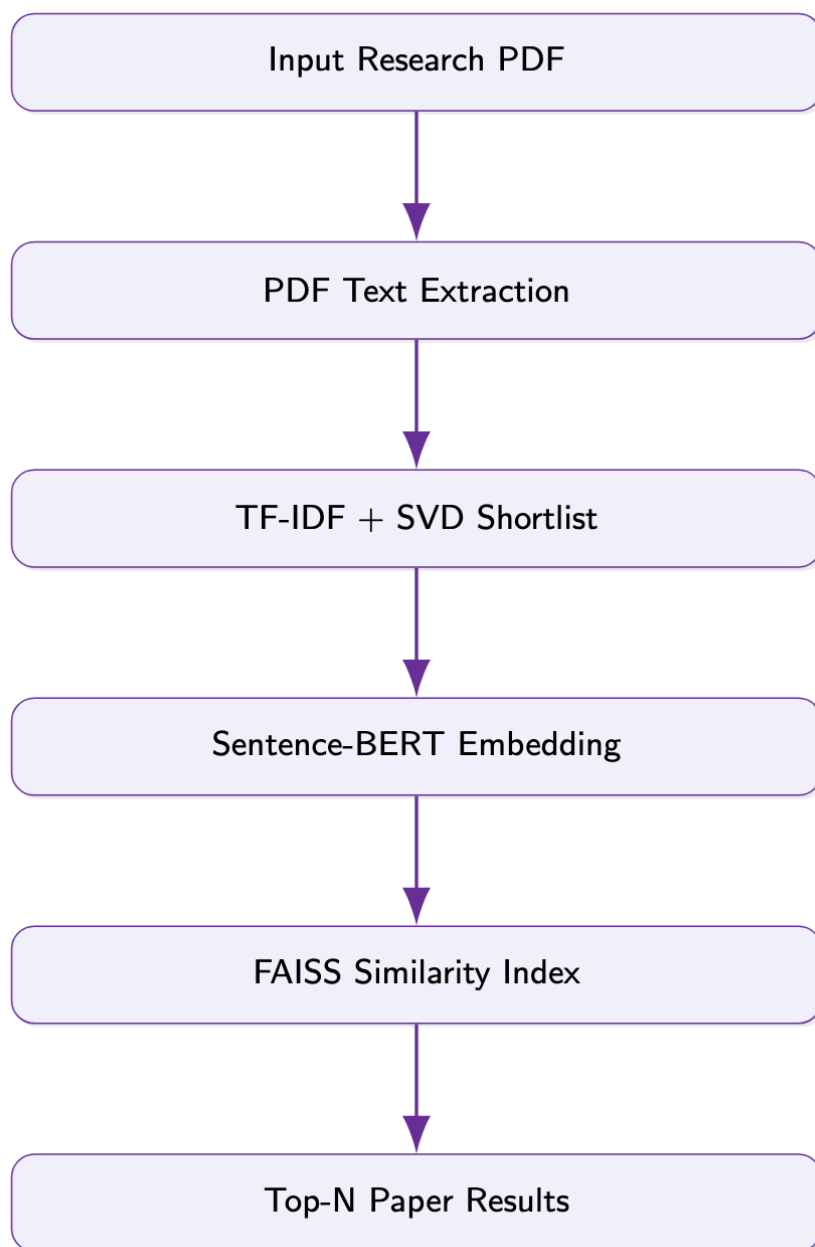
5. System Architecture

The architecture integrates multiple NLP and machine learning components efficiently.

Workflow:

1. Input PDF → Text Extraction
2. Text Vectorization (TF-IDF + SVD)
3. Candidate Shortlisting
4. Sentence Embedding Generation
5. FAISS Similarity Search
6. Ranked Paper Recommendations

Architecture Diagram (Conceptual)



6. Conclusion and Future Work

The NLP-based hybrid recommendation system effectively identifies semantically similar research papers using a combination of SVD-based LSA and Sentence-BERT embeddings. It successfully integrates traditional IR and modern deep learning to deliver context-aware recommendations.

Key Outcomes

- Achieved fast and accurate retrieval for large datasets.
- Demonstrated end-to-end integration in Google Colab.
- Validated the pipeline with real arXiv data.

Future Enhancements

- Add citation networks and author-based similarity.
- Integrate a web-based UI for real-time use.
- Implement topic modeling (LDA/BERTTopic) for thematic clustering.
- Extend to multilingual academic papers.

References

1. Devlin, J., et al. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. NAACL-HLT.
2. Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. EMNLP.
3. Facebook AI Research. FAISS: A library for efficient similarity search and clustering of dense vectors.
4. Cornell University arXiv Dataset (Kaggle).
5. Scikit-learn Documentation (TF-IDF, SVD Modules).